

Creating AutoML with Azure Python SDK

Reference :

1. [\(https://github.com/binodsuman/azure_ml\)](https://github.com/binodsuman/azure_ml) (OG Content)
2. [\(https://learn.microsoft.com/en-us/training/modules/intro-to-azure-machine-learning-service/4-azure-ml-experiment\)](https://learn.microsoft.com/en-us/training/modules/intro-to-azure-machine-learning-service/4-azure-ml-experiment)
3. [\(https://www.youtube.com/watch?v=Lxdui2F2Zw8&ab_channel=AmbarishGanguly\)](https://www.youtube.com/watch?v=Lxdui2F2Zw8&ab_channel=AmbarishGanguly) (Datastore)
4. [\(https://github.com/anikch/azure-ml\)](https://github.com/anikch/azure-ml)
5. MLOPS : [\(https://www.sqlservergeeks.com/azure-machine-learning-for-mlops-by-mercy-ranjith/\)](https://www.sqlservergeeks.com/azure-machine-learning-for-mlops-by-mercy-ranjith/)

STEP 1: Go to azure and create an account from [\(https://portal.azure.com/#home\)](https://portal.azure.com/#home)

STEP 2 : Go to Anaconda and create an environment to work with azure ML specifically.

Follow the steps to create a virtual env in Azure:

- Create an environment : conda create --name azure_ml_env
- Activate the virtual env : conda activate azure_ml_env

STEP 3: Install python azure-ml sdk : pip install azureml-sdk

This step takes time and if some of the installation fails, do upgrade pip version by the command : python -m pip install --upgrade pip

In [1]: pip show azureml-sdk ### To show azureml-sdk version

```
Name: azureml-sdk
Version: 1.49.0
Summary: Used to build and run machine learning workflows upon the Azure Machine Learning service.
Home-page: https://docs.microsoft.com/python/api/overview/azure/ml/?view=azure-ml-py (https://docs.microsoft.com/python/api/overview/azure/ml/?view=azure-ml-py)
Author: Microsoft Corp
Author-email:
License: https://aka.ms/azureml-sdk-license (https://aka.ms/azureml-sdk-license)
Location: c:\users\krish\anaconda3\envs\azure_ml_env\lib\site-packages
Requires: azureml-core, azureml-dataset-runtime, azureml-pipeline, azureml-train-automl-client, azureml-train-core
Required-by:
---
Name: azureml-sdk
Version: 1.49.0
Summary: Used to build and run machine learning workflows upon the Azure Machine Learning service.
Home-page: https://docs.microsoft.com/python/api/overview/azure/ml/?view=azure-ml-py (https://docs.microsoft.com/python/api/overview/azure/ml/?view=azure-ml-py)
Author: Microsoft Corp
Author-email:
License: https://aka.ms/azureml-sdk-license (https://aka.ms/azureml-sdk-license)
Location: c:\users\krish\anaconda3\envs\azure_ml_env\lib\site-packages
Requires: azureml-core, azureml-dataset-runtime, azureml-pipeline, azureml-train-automl-client, azureml-train-core
Required-by:
Note: you may need to restart the kernel to use updated packages.

WARNING: Package(s) not found: ###, To, show, version
```

STEP 4: Install Jupyter by using : pip install jupyter

STEP 5: Launch jupyter notebook by using command: jupyter notebook

STEP 6: Go to azure -> Subscription and check the subscription id you have.

STEP 7: Importing the necessary libraries

In [2]: import warnings
warnings.filterwarnings('ignore')

```
In [3]: import azureml
import azure.core
from azureml.core import Workspace, Environment, Experiment, ScriptRunConfig
from azureml.core.compute import AmlCompute, ComputeTarget
from azureml.core.compute_target import ComputeTargetException
from azureml.core import Model
```

STEP 8: Specifying the suitable configuration info

```
In [4]: resource_name = "KD_resource_group"
workspace_name = "KD_demo_workspace"
subscriptionID = "da7f92e5-f637-4517-90fb-493d47170db2" # Please enter your
aml_compute_target = "KD-cluster" # ALL SMALL LETTER,
experiment_name= 'demo_expirement_1'
environment_name = 'azure_ml_demo'
```

STEP 9: Creating a new workspace if not created

```
In [5]: # Now create Workspace if no workspace is there
try:
    ws=Workspace.from_config()
    print('Workspace is already exist')
except:
    ws=Workspace.create(workspace_name,
                         resource_group=resource_name,
                         create_resource_group=False, # If this is made true
                         subscription_id=subscriptionID,
                         location="Central India")
    ws.write_config('.azureml')
```

Performing interactive authentication. Please follow the instructions on the terminal.

The default web browser has been opened at <https://login.microsoftonline.com/organizations/oauth2/v2.0/authorize>. (<https://login.microsoftonline.com/organizations/oauth2/v2.0/authorize>.) Please continue the login in the web browser. If no web browser is available or if the web browser fails to open, use device code flow with `az login --use-device-code`.

Interactive authentication successfully completed.
Deploying StorageAccount with name kdodemowostorage3c2969387.
Deploying AppInsights with name kdodemowoinsights50a9da65.
Deployed AppInsights with name kdodemowoinsights50a9da65. Took 8.69 seconds.
Deploying Workspace with name KD_demo_workspace.
Deployed Workspace with name KD_demo_workspace. Took 37.17 seconds.

STEP 10: After the above cell got executed, cross check from the Azure portal. Go to Azure -> home -> resource group(under which the workspace is supposed to get created) -> check if the workspace has got created or not.

The screenshot shows the Azure portal interface for a resource group named 'KD_resource_group'. The left sidebar contains navigation links for Overview, Activity log, Access control (IAM), Tags, Resource visualizer, Events, Settings (Deployments, Security, Stacks, Policies, Properties, Locks), and a search bar. The main content area displays the 'Essentials' section with details like Subscription (move) to 'Free Trial', Subscription ID 'da7f92e5-f637-4517-90fb-493d47170db2', and Location 'Central India'. Below this is a table titled 'Resources' showing four items: 'KD_demo_workspace' (Azure Machine Learning workspace), 'kddemowoinights50a9da65' (Application Insights), 'kddemowokeyvaultd46ac110' (Key vault), and 'kddemowostorage3c2969387' (Storage account). A filter bar at the top allows filtering by Name, Type, and Location.

```
In [5]: ws=Workspace.from_config() ## To fetch workspace
ws
```

```
Out[5]: Workspace.create(name='KD_demo_workspace', subscription_id='da7f92e5-f637-4517-90fb-493d47170db2', resource_group='KD_resource_group')
```

```
In [6]: print("SDK version:", azureml.core.VERSION)
print(ws.name, ws.resource_group, ws.location, ws.subscription_id, sep = '\n')
```

```
SDK version: 1.49.0
KD_demo_workspace
KD_resource_group
centralindia
da7f92e5-f637-4517-90fb-493d47170db2
```

STEP 11: Now, we need to click on the "Launch Studio" button or the 'Studio web url' to go to the Azure MI Studio.

Link for our case : [https://ml.azure.com/?tid=d8533050-79a6-4371-8c5b-ddfad9be8d2&wsid=/subscriptions/b194183f-1d92-4868-8e02-c46f33bc7030/resourcegroups/KD_resource_group/providers/Microsoft.MachineLearningServices](https://ml.azure.com/?tid=d8533050-79a6-4371-8c5b-ddfad9be8d2&wsid=/subscriptions/b194183f-1d92-4868-8e02-c46f33bc7030/resourcegroups/KD_resource_group/providers/Microsoft.MachineLearningServices((https://ml.azure.com/?tid=d8533050-79a6-4371-8c5b-ddfad9be8d2&wsid=/subscriptions/b194183f-1d92-4868-8e02-c46f33bc7030/resourcegroups/KD_resource_group/providers/Microsoft.MachineLearningServices)

The below screen must come up if all the things are done properly as of now.

The screenshot shows the Azure portal interface for an Azure Machine Learning workspace named 'KD_demo_workspace'. The left sidebar contains Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, and Events. The main content area displays the 'Essentials' section with details like Resource group 'KD_resource_group', Location 'Central India', Subscription 'Free Trial', Subscription ID 'da7f92e5-f637-4517-90fb-493d47170db2', and Storage 'kddemowostorage3c2969387'. It also lists 'Studio web URL' as 'https://ml.azure.com/?tid=9987f6ff-f026-43d3-9042-403c88d...', 'Container Registry' as '...', 'Key Vault' as 'kddemowokeyvaultd46ac110', 'Application Insights' as 'kddemowoinights50a9da65', and 'MLflow tracking URI' as 'azurerm://centralindia.api.azureml.ms/mlflow/v1.0/subscriptio...'. A large gray button at the bottom is partially visible.

STEP 12: Check compute tab at the left of the screen and see if there is any compute instance already created or not. If not then we would have to create a new compute instance to run our ML model. The following code will create a compute instance.

```
In [8]: # Create Compute Target if it's not created already
try:
    aml_compute = AmlCompute(ws, aml_compute_target)
    print("This Compute Target already exist.")
except ComputeTargetException:
    print("creating new compute target : ",aml_compute_target)
    provisioning_config = AmlCompute.provisioning_configuration(vm_size = "S
                                min_nodes =
                                max_nodes =
                                idle_seconds_before_scaledown
    aml_compute = ComputeTarget.create(ws, aml_compute_target, provisioning_
    aml_compute.wait_for_completion(show_output=True, min_node_count=None, t
print("Azure Machine Learning Compute attached now")
```

```
creating new compute target : KD-cluster
InProgress.....
SucceededProvisioning operation finished, operation "Succeeded"
Succeeded.....
AmlCompute wait for completion finished

Minimum number of nodes requested have been provisioned
Azure Machine Learning Compute attached now
```

STEP 13 After successful execution of the command, we should see the compute instance getting attached in the compute tab. path : azure -> home -> resource_group -> KD_resource_group -> demo_azure_ml_workspace -> Launch studio -> compute -> compute clusters.

Cluster node status		Cluster state
 1	Allocation state ● Succeeded (1 node) Allocation state transition time 7/8/2023, 1:37:10 PM Created on 7/8/2023, 1:35:28 PM Current node count 1	
Attributes <ul style="list-style-type: none"> Compute name: KD-cluster Resource ID: ... Compute type: Machine Learning compute Subscription ID: da792e5-1637-4517-90fb-493d47170db2 Resource group: KD_resource_group Workspace: KD_demo_workspace Region: centralindia 	Resource properties <ul style="list-style-type: none"> Virtual machine size: Standard_D2_v2 (2 cores, 7 GB RAM, 100 GB disk) Processing unit: CPU - General purpose Estimated cost: \$0.17/hr per node OS Type: Linux Virtual machine tier: Dedicated Minimum number of nodes: 1 Maximum number of nodes: 4 Idle seconds before scale down: 3000 SSH access: Disabled Virtual network/subnet: ... 	

```
In [7]: ## Checking the compute target that has been created
for compute_name in ws.compute_targets:
    compute = ws.compute_targets[compute_name]
    print(compute.name, ":", compute.type)
```

KD-cluster : AmlCompute

STEP 14: Create Experiment using the following command and check if the experiment has been created successfully or not from the path azure -> home -> resource_group -> KD_resource_group -> demo_azure_ml_workspace -> Launch studio -> jobs

```
In [13]: exp=Experiment(workspace=ws, name=experiment_name,
run = exp.start_logging()
run.complete() # Start Logging data fr
# Complete the experime
```

In [11]: exp

Out[11]:

Name	Workspace	Report Page	Docs Page
demo_expiration_1	KD_demo_workspace	Link to Azure Machine Learning studio (https://ml.azure.com/experiments/fce22d92-4d13-4618-9429-71a79d6b3f01?wsid=subscriptions/da7f92e5-f637-4517-90fb-493d47170db2/resourcegroups/KD_resource_group/wf02643d3-9042-403c88da284a)	Link to Documentation (https://docs.microsoft.com/en-us/python/api/azureml-core/azureml.core.experiment.Experiment?view=azure-ml-py)

Display name	Status	Created on	Duration	Created by	Compute target	Job type	Tags
coral_camera_k873b9zd	Completed	Jul 8, 2023 2:30 PM	8s	Krishnendu Dey	local		

Logging

In [39]:

```
import pandas as pd
import seaborn as sns
from sklearn.datasets import load_diabetes
```

```
data = load_diabetes()
X = pd.DataFrame(data['data'], columns=data.feature_names)
Y = pd.Series(data['target'])
row_count = len(X)
```

In [40]:

```
# Log the row count
run.log('observations', row_count)
run.complete()
```

Retrieving the logs

In [41]: *### The Logs can be viewed using RunDetails*

```
from azureml.widgets import RunDetails
RunDetails(run).show()
```

```
_UserRunWidget(widget_settings={'childWidgetDisplay': 'popup', 'send_telemetry': False, 'log_level': 'INFO', '...
```

In [43]: *import json*

```
# Get logged metrics
metrics = run.get_metrics()
print(json.dumps(metrics, indent=2))
```

```
{
    "observations": [
        8,
        442
    ]
}
```

STEP 15: Working with datastore:

In [72]: *### Fetching the default datastores*

```
default_ds=ws.get_default_datastore()
default_ds
```

Out[72]: {

```
"name": "workspaceblobstore",
"container_name": "azureml-blobstore-8f813250-e78c-4309-9f63-32f488470208",
"account_name": "kddemowostorage3c2969387",
"protocol": "https",
"endpoint": "core.windows.net"
}
```

Name	Type	Storage name	Created on
azureml	Azure Blob Storage	kddemowostorage3c2969387	Jul 8, 2023 5:30 PM
workspacefilestore	Azure file share	kddemowostorage3c2969387	Jul 8, 2023 1:26 PM
workspaceartifactstore	Azure Blob Storage	kddemowostorage3c2969387	Jul 8, 2023 1:26 PM
workspaceworkingdirectory	Azure file share	kddemowostorage3c2969387	Jul 8, 2023 1:26 PM
workspaceblobstore (Default)	Azure Blob Storage	kddemowostorage3c2969387	Jul 8, 2023 1:26 PM

In [73]: *### Uploading a file in the default ds*

```
default_ds.upload_files(files=['diabetes.csv'],
                        target_path='diabetes_data/',
                        overwrite=True,
                        show_progress=True)
```

"datastore.upload_files" is deprecated after version 1.0.69. Please use "FileDatasetFactory.upload_directory" instead. See Dataset API change notice at <https://aka.ms/dataset-deprecation>. (<https://aka.ms/dataset-deprecation>.)

Uploading an estimated of 1 files

Uploading diabetes.csv

Uploaded diabetes.csv, 1 files out of an estimated total of 1

Uploaded 1 files

Out[73]: \$AZUREML_DATAREFERENCE_789addfab01e440b852a7af2a17b187f

The screenshot shows the Azure AI | Machine Learning Studio interface. On the left, there's a sidebar with navigation links like Home, Model catalog, Notebooks, Automated ML, Designer, Prompt flow, Data, Jobs, Components, Pipelines, Environments, Models, Endpoints, Compute, Monitoring, Data Labeling, and Linked Services. The 'Data' link is currently selected. In the main area, the path 'Default Directory > KD_demo_workspace > Data > workspaceblobstore' is shown. The title 'workspaceblobstore (Default)' is displayed with a star icon. Below it, there are two tabs: 'Overview' (selected) and 'Browse'. Under 'Overview', there are sections for General, Data assets, and Metrics. The General section shows details such as Datastore name (workspaceblobstore), Datastore type (Azure Blob Storage), Created by (Service Principal), Subscription ID (da792e5-f637-4517-90fb-493d47170db2), Resource group name (KD_resource_group), Protocol (https), Endpoint (core.windows.net), Account name (kddemowostorage3c2969387), Blob container (azureml-blobstore-8f813250-e78c-4309-9f63-32f488470208), and Storage URI (https://kddemowostorage3c2969387.blob.core.windows.net/azureml-blobstore-8f813250-e78c-4309-9f63-32f488470208). The Data assets section shows a table titled 'Browse preview' with columns Name, Created on, and Modified on. It lists two items: 'azureml' and 'diabetes_data', both created and modified on Jul 8, 2023.

Name	Created on	Modified on
azureml	--	--
diabetes_data	--	--

```
In [74]: #### Reading data in the form of dataset from datastore
from azureml.core import Dataset
dataset=Dataset.Tabular.from_delimited_files(path=(default_ds, 'diabetes_data.csv'))
df=dataset.to_pandas_dataframe()
df.head(5)
```

Out[74]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.628
1	1	85	66	29	0	26.6	0.339
2	8	183	64	0	0	23.3	0.601
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.287



```
In [75]: ### Registering the dataset
try:
    dataset = dataset.register(workspace=ws,
                               name='diabetes dataset',
                               description='diabetes data',
                               create_new_version=True,
                               tags={'format':'csv'})
except Exception as ex :
    print(ex)
```

```
In [76]: ### Fetching all the datafiles
datafiles=Dataset.File.from_files(path=(default_ds, 'diabetes_data/diabetes.csv'))
for filepath in datafiles.to_path():
    print(filepath)
```

/diabetes.csv

```
In [77]: ### Registering datafiles
try:
    datafiles=datafiles.register(workspace=ws,
                                 name='diabetes files',
                                 description='diabetes files',
                                 create_new_version=True,
                                 tags={'format':'csv'})
except Exception as ex :
    print(ex)
```

```
In [83]: # print(dataset)
# print(datafiles)
```

```
In [82]: ## Fetching the diabetes dataset
diabetes_ds=ws.datasets.get('diabetes dataset')
# Alternate Query: Dataset.get_by_name(workspace=ws, name='diabetes dataset')
diabetes_ds
```

```
Out[82]: {
    "source": [
        "('workspaceblobstore', 'diabetes_data/diabetes.csv')"
    ],
    "definition": [
        "GetDatastoreFiles",
        "ParseDelimited",
        "DropColumns",
        "SetColumnTypes"
    ],
    "registration": {
        "id": "ba38ecfd-9fd7-45b1-830b-888e2b03dd7d",
        "name": "diabetes dataset",
        "version": 1,
        "description": "diabetes data",
        "tags": {
            "format": "csv"
        },
        "workspace": "Workspace.create(name='KD_demo_workspace', subscription_id='da7f92e5-f637-4517-90fb-493d47170db2', resource_group='KD_resource_group')"
    }
}
```

```
In [84]: input_data = diabetes_ds.as_named_input('dataset')
input_data
```

```
Out[84]: <azureml.data.dataset_consumption_config.DatasetConsumptionConfig at 0x2788 da8f6a0>
```

STEP 15: Creating environment and training model:

In this step we actually specify our required libraries and packages that are to be needed to run the code in the .py file. That environment will be set up in the azure VMs and our code will be executed through the vm in the cloud. To run a script as an experiment, you must define a script configuration that defines the script to be run and the Python environment in which to run it. This is implemented by using a ScriptRunConfig object. This step takes approx 15 mins to execute.

```
In [8]: # Creating environment : make sure you are in the same directory where mytra  
env = Environment.from_conda_specification(name=environment_name,  
                                              file_path="./envfile.yml",)  
env.register(workspace=ws) # Re  
env
```

Out[8]: {

```
    "assetId": null,
    "databricks": {
        "eggLibraries": [],
        "jarLibraries": [],
        "mavenLibraries": [],
        "pypiLibraries": [],
        "rcranLibraries": []
    },
    "docker": {
        "arguments": [],
        "baseDockerfile": null,
        "baseImage": "mcr.microsoft.com/azureml/openmpi4.1.0-ubuntu20.04:20
230120.v1",
        "baseImageRegistry": {
            "address": null,
            "password": null,
            "registryIdentity": null,
            "username": null
        },
        "buildContext": null,
        "enabled": false,
        "platform": {
            "architecture": "amd64",
            "os": "Linux"
        },
        "sharedVolumes": true,
        "shmSize": "2g"
    },
    "environmentVariables": {
        "EXAMPLE_ENV_VAR": "EXAMPLE_VALUE"
    },
    "inferencingStackVersion": null,
    "name": "azure_ml_demo",
    "python": {
        "baseCondaEnvironment": null,
        "condaDependencies": {
            "channels": [
                "defaults"
            ],
            "dependencies": [
                "python=3.10.9",
                "anaconda",
                "pip",
                {
                    "pip": [
                        "azureml-sdk",
                        "mlflow",
                        "azureml-mlflow"
                    ]
                }
            ],
            "name": "azure_ml"
        },
        "condaDependenciesFile": null,
        "interpreterPath": "python",
        "userManagedDependencies": false
    }
}
```

```
},
"r": null,
"spark": {
    "packages": [],
    "precachePackages": true,
    "repositories": []
},
"version": null
}
```

Note : Changing hyperparameters and training the model and validating it again by triggering a new job altogether. To use parameters in a script, you must use a library such as **argparse** in the mytrain.py file to read the arguments passed to the script and assign them to variables

```
In [10]: ## ScriptRunConfig is used to run the machine learning model from notebook w
config=ScriptRunConfig(source_directory=". /", script="mytrain.py",
                      compute_target=aml_compute_target,
                      environment=env,)

execution=exp.submit(config)
execution.wait_for_completion(show_output=True)
```

```
RunId: demo_expirement_1_1688811327_cb3ddc60
Web View: https://ml.azure.com/runs/demo\_expirement\_1\_1688811327\_cb3ddc60?w
sid=/subscriptions/da7f92e5-f637-4517-90fb-493d47170db2/resourcegroups/KD_r
esource_group/workspaces/KD_demo_workspace&tid=9987f6ff-f026-43d3-9042-403c
88da284a (https://ml.azure.com/runs/demo\_expirement\_1\_1688811327\_cb3ddc60?w
sid=/subscriptions/da7f92e5-f637-4517-90fb-493d47170db2/resourcegroups/KD_r
esource_group/workspaces/KD_demo_workspace&tid=9987f6ff-f026-43d3-9042-403c
88da284a)
```

```
Streaming user_logs/std_log.txt
=====
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
Accuracy on training set: 1.000
Accuracy on test set: 0.766
Accuracy 0.7662337662337663
0.7662337662337663
Cleaning up all outstanding Run operations, waiting 300.0 seconds
1 items cleaning up...
Cleanup took 0.10429525375366211 seconds
```

```
Execution Summary
=====
```

```
RunId: demo_expirement_1_1688811327_cb3ddc60
Web View: https://ml.azure.com/runs/demo\_expirement\_1\_1688811327\_cb3ddc60?w
sid=/subscriptions/da7f92e5-f637-4517-90fb-493d47170db2/resourcegroups/KD_r
esource_group/workspaces/KD_demo_workspace&tid=9987f6ff-f026-43d3-9042-403c
88da284a (https://ml.azure.com/runs/demo\_expirement\_1\_1688811327\_cb3ddc60?w
sid=/subscriptions/da7f92e5-f637-4517-90fb-493d47170db2/resourcegroups/KD_r
esource_group/workspaces/KD_demo_workspace&tid=9987f6ff-f026-43d3-9042-403c
88da284a)
```

```
Out[10]: {'runId': 'demo_expirement_1_1688811327_cb3ddc60',
  'target': 'KD-cluster',
  'status': 'Completed',
  'startTimeUtc': '2023-07-08T10:15:43.09724Z',
  'endTimeUtc': '2023-07-08T10:15:57.475597Z',
  'services': {},
  'properties': {'_azureml.ComputeTargetType': 'amlctrain',
    'ContentSnapshotId': '1aac51a8-6f10-4878-9283-68729dbb58ba',
    'ProcessInfoFile': 'azureml-logs/process_info.json',
    'ProcessStatusFile': 'azureml-logs/process_status.json'},
  'inputDatasets': [],
  'outputDatasets': [],
  'runDefinition': {'script': 'mytrain.py',
    'command': '',
    'useAbsolutePath': False,
    'arguments': ['--min_samples_leaf', '5', '--min_samples_split', '7'],
    'sourceDirectoryDataStore': None,
    'framework': 'Python',
    'communicator': 'None',
    'target': 'KD-cluster',
    'dataReferences': {},
    'data': {},
    'outputData': {},
    'datacaches': [],
    'jobName': None,
    'maxRunDurationSeconds': 2592000,
    'nodeCount': 1,
    'instanceTypes': [],
    'priority': None,
    'credentialPassthrough': False,
    'identity': None,
    'environment': {'name': 'azure_ml_demo',
      'version': '1',
      'assetId': 'azureml://locations/centralindia/workspaces/8f813250-e78c-4309-9f63-32f488470208/environments/azure_ml_demo/versions/1',
      'autoRebuild': True,
      'python': {'interpreterPath': 'python',
        'userManagedDependencies': False,
        'condaDependencies': {'name': 'azure_ml',
          'channels': ['defaults'],
          'dependencies': ['python=3.10.9',
            'anaconda',
            'pip',
            {'pip': ['azureml-sdk', 'mlflow', 'azureml-mlflow']}]}},
      'baseCondaEnvironment': None},
    'environmentVariables': {'EXAMPLE_ENV_VAR': 'EXAMPLE_VALUE'},
    'docker': {'baseImage': 'mcr.microsoft.com/azureml/openmpi4.1.0-ubuntu20.04:20230120.v1',
      'platform': {'os': 'Linux', 'architecture': 'amd64'},
      'baseDockerfile': None,
      'baseImageRegistry': {'address': None, 'username': None, 'password': None},
      'enabled': False,
      'arguments': []},
    'spark': {'repositories': [], 'packages': [], 'precachePackages': True},
    'inferencingStackVersion': None},
    'history': {'outputCollection': True},
```

```
'directoriesToWatch': ['logs'],
'enableMLflowTracking': True,
'snapshotProject': True},
'spark': {'configuration': {'spark.app.name': 'Azure ML Experiment',
'spark.yarn.maxAppAttempts': '1'}},
'parallelTask': {'maxRetriesPerWorker': 0,
'workerCountPerNode': 1,
'terminalExitCodes': None,
'configuration': {}},
'amlCompute': {'name': None,
'vemSize': None,
'retainCluster': False,
'clusterMaxNodeCount': None},
'aiSuperComputer': {'instanceType': 'D2',
'imageVersion': None,
'location': None,
'aiSuperComputerStorageData': None,
'interactive': False,
'scalePolicy': None,
'verticalClusterArmId': None,
'tensorboardLogDirectory': None,
'sshPublicKey': None,
'sshPublicKeys': None,
'enableAzmlInt': True,
'priority': 'Medium',
'slaTier': 'Standard',
'userAlias': None},
'kubernetesCompute': {'instanceType': None},
'tensorflow': {'workerCount': 1, 'parameterServerCount': 1},
'mpi': {'processCountPerNode': 1},
'pyTorch': {'communicationBackend': 'nccl', 'processCount': None},
'hdi': {'yarnDeployMode': 'Cluster'},
'containerInstance': {'region': None, 'cpuCores': 2.0, 'memoryGb': 3.5},
'exposedPorts': None,
'docker': {'useDocker': False,
'sharedVolumes': True,
'shmSize': '2g',
'arguments': []},
'cmk8sCompute': {'configuration': {}},
'commandReturnCodeConfig': {'returnCode': 'Zero',
'successfulReturnCodes': []},
'environmentVariables': {},
'applicationEndpoints': {},
'parameters': [],
'logFiles': {'user_logs/std_log.txt': 'https://kddemowostorage3c2969387.blob.core.windows.net/azureml/ExperimentRun/dcid.demo_expirement_1_1688811327_cb3ddc60/user_logs/std_log.txt?sv=2019-07-07&sr=b&sig=IqvdwYW4%2B1BErti6J%2FP2sd84l4L7N6lgGGg0XKzAYYg%3D&skoid=519c9b80-fd48-42e8-a412-61c41ef4a583&ktid=9987f6ff-f026-43d3-9042-403c88da284a&skt=2023-07-08T09%3A32%3A06Z&ske=2023-07-09T17%3A42%3A06Z&skv=2019-07-07&st=2023-07-08T10%3A06%3A00Z&se=2023-07-08T18%3A16%3A00Z&sp=r',
'system_logs/cs_capability/cs-capability.log': 'https://kddemowostorage3c2969387.blob.core.windows.net/azureml/ExperimentRun/dcid.demo_expirement_1_1688811327_cb3ddc60/system_logs/cs_capability/cs-capability.log?sv=2019-07-07&sr=b&sig=sShAt3wxIL9jo0rUbHMxhSDnu%2FkFTUGnhJaIx58N2zs%3D&skoid=519c9b80-fd48-42e8-a412-61c41ef4a583&skt=2023-07-08T09%3A32%3A06Z&ske=2023-07-09T17%3A42%3A06Z&skv=2019-07-07'}
```

```
&st=2023-07-08T10%3A06%3A01Z&se=2023-07-08T18%3A16%3A01Z&sp=r',
    'system_logs/hosttools_capability/hosttools-capability.log': 'https://kddemowostorage3c2969387.blob.core.windows.net/azureml/ExperimentRun/dcida.demo_expirement_1_1688811327_cb3ddc60/system_logs/hosttools_capability/hosttools-capability.log?sv=2019-07-07&sr=b&sig=ZVGoPjmFH3VGau9xtK4QRuawMH0CG8gNWNUokaERGRE%3D&skoid=519c9b80-fd48-42e8-a412-61c41ef4a583&sktid=9987f6ff-f026-43d3-9042-403c88da284a&skt=2023-07-08T09%3A32%3A06Z&ske=2023-07-09T17%3A42%3A06Z&sks=b&skv=2019-07-07&st=2023-07-08T10%3A06%3A01Z&se=2023-07-08T18%3A16%3A01Z&sp=r',
    'system_logs/lifecycler/execution-wrapper.log': 'https://kddemowostorage3c2969387.blob.core.windows.net/azureml/ExperimentRun/dcida.demo_expirement_1_1688811327_cb3ddc60/system_logs/lifecycler/execution-wrapper.log?sv=2019-07-07&sr=b&sig=4EY9pgFoALfDmIU%2Blu%2FewwfUTUKsHysIoMOTAUX6lg%3D&skoid=519c9b80-fd48-42e8-a412-61c41ef4a583&sktid=9987f6ff-f026-43d3-9042-403c88da284a&skt=2023-07-08T09%3A32%3A06Z&ske=2023-07-09T17%3A42%3A06Z&sks=b&skv=2019-07-07&st=2023-07-08T10%3A06%3A01Z&se=2023-07-08T18%3A16%3A01Z&sp=r',
    'system_logs/lifecycler/lifecycler.log': 'https://kddemowostorage3c2969387.blob.core.windows.net/azureml/ExperimentRun/dcida.demo_expirement_1_1688811327_cb3ddc60/system_logs/lifecycler/lifecycler.log?sv=2019-07-07&sr=b&sig=JV2PedtG%2FrCz4Tg315de61ceitDEY%2FoieO64C5wXLSY%3D&skoid=519c9b80-fd48-42e8-a412-61c41ef4a583&sktid=9987f6ff-f026-43d3-9042-403c88da284a&skt=2023-07-08T09%3A32%3A06Z&ske=2023-07-09T17%3A42%3A06Z&sks=b&skv=2019-07-07&st=2023-07-08T10%3A06%3A01Z&se=2023-07-08T18%3A16%3A01Z&sp=r',
    'system_logs/metrics_capability/metrics-capability.log': 'https://kddemowostorage3c2969387.blob.core.windows.net/azureml/ExperimentRun/dcida.demo_expirement_1_1688811327_cb3ddc60/system_logs/metrics_capability/metrics-capability.log?sv=2019-07-07&sr=b&sig=0jen4RTXPloeYkt0aC%2BlkxUvNzi43ok1E%2BTBeJqDH3U%3D&skoid=519c9b80-fd48-42e8-a412-61c41ef4a583&sktid=9987f6ff-f026-43d3-9042-403c88da284a&skt=2023-07-08T09%3A32%3A06Z&ske=2023-07-09T17%3A42%3A06Z&sks=b&skv=2019-07-07&st=2023-07-08T10%3A06%3A01Z&se=2023-07-08T18%3A16%3A01Z&sp=r',
    'system_logs/snapshot_capability/snapshot-capability.log': 'https://kddemowostorage3c2969387.blob.core.windows.net/azureml/ExperimentRun/dcida.demo_expirement_1_1688811327_cb3ddc60/system_logs/snapshot_capability/snapshot-capability.log?sv=2019-07-07&sr=b&sig=enxAKpFOM7nWAjI%2BxdfNnHjHZzPJkz6HY06dCS1Pzhw%3D&skoid=519c9b80-fd48-42e8-a412-61c41ef4a583&sktid=9987f6ff-f026-43d3-9042-403c88da284a&skt=2023-07-08T09%3A32%3A06Z&ske=2023-07-09T17%3A42%3A06Z&sks=b&skv=2019-07-07&st=2023-07-08T10%3A06%3A01Z&se=2023-07-08T18%3A16%3A01Z&sp=r'},
    'submittedBy': 'Krishnendu Dey'}
```

When the code gets successfully executed, then go to jobs and check if the job id has completed or not from azure.

The screenshot shows the Azure AI | Machine Learning Studio interface. A job named 'teal_goat_t0bbbv5d' is displayed, showing it is completed. The 'Outputs + logs' tab is selected, indicating the model has been run.

STEP 16: To check how the model performs, go to outputs and logs tab.

The screenshot shows the Azure AI | Machine Learning Studio interface. A job named 'affable_oxygen_8rgvws3s' is displayed, showing it is completed. The 'Outputs + logs' tab is selected, and the 'std_log.txt' file is selected in the log viewer, displaying its contents.

In [15]: execution

Out[15]:	Experiment	Id	Type	Status	Details Page	Docs Page
	demo_expirement1	demo_expirement1_u1n888d1p27ncb3dd60	Completed	Link to Azure Machine Learning studio (https://ml.azure.com/documents/tutorials wsid=/subscriptions/31da3f92e55.microsoft.com/resourceGroups/f637-4517-us/python/api/azureml.core/f93d47170db2/resources/mD_f026-43d3-9042-403c88da284a)	Link to Azure Machine Learning studio (https://ml.azure.com/documents/tutorials wsid=/subscriptions/31da3f92e55.microsoft.com/resourceGroups/f637-4517-us/python/api/azureml.core/f93d47170db2/resources/mD_f026-43d3-9042-403c88da284a)	Link to https://ml.azure.com/documents/tutorials wsid=/subscriptions/31da3f92e55.microsoft.com/resourceGroups/f637-4517-us/python/api/azureml.core/f93d47170db2/resources/mD_f026-43d3-9042-403c88da284a)

STEP 17: Now the above steps would not give us anything getting displayed on the Metrics tab for the job. To populate the results of the population metrics on the tab we need to make certain modification to the code. The code with specific changes can be found in the mytrain_log.py file.

```
In [89]: config=ScriptRunConfig(source_directory=".",
                             script="train.py",
                             compute_target=aml_compute_target,
                             environment=env,
                             arguments=['--min_samples_leaf',5,
                                         '--min_samples_split',7,
                                         '--input_data',input_data,])
execution=exp.submit(config)
```

```
In [90]: execution
```

Out[90]:	Experiment	Id	Type	Status	Details Page	Docs Page
	demo_expirement	demo_expirement_1_1688823904_n62bcc26	Starting		Link to Azure Machine Learning studio (https://ml.azure.com/demos/expirement_1_1688823904_n62bcc26/f637-4517-90fb-493d47170db2/resourcegroup/KD_f026-43d3-9042-403c88da284a)	Link to Document

The screenshot shows the Azure AI | Machine Learning Studio interface. On the left, there's a sidebar with navigation links: All workspaces, Home, Model catalog (PREVIEW), Authoring (Notebooks, Automated ML, Designer, Prompt flow PREVIEW), Assets (Data, Jobs, Components, Pipelines, Environments). The main area shows a job named 'yellow_guitar_sj6fc5n6' under the 'Jobs' section. A message indicates it's using the new compute runtime. The job status is 'Completed'. Below the job name, there are tabs for Overview, Metrics, Images, Child jobs, Outputs + logs, Code, Explanations (preview), Fairness (preview), and Monitoring. Under the Metrics tab, there's a chart and some numerical values. A modal window titled 'Select metrics' is open, showing 'Accuracy' with a value of '0.7662338' and 'Parameter' with a long JSON string.

STEP 19: So far the model has been built in the backend and the training and testing is done using azure VM. But if we want to see the model from the UI , we won't be able to coz the model would not be there in the models option of the Azure ML Studio untill we register it for further usage. Below code demonstrates how the model registration is done.

```
In [69]: #Register model in workspace
model_name = "diabetic_model_1"
model = Model.register(ws,
                      model_path=".outputs/diabeticmodel.pkl",
                      model_name=model_name,
                      model_framework=Model.Framework.SCIKITLEARN,
                      description='A classification model',
                      tags={'data-format': 'CSV'},)
```

Registering model diabetic_model_1

```
In [70]: model
```

```
Out[70]: Model(workspace=Workspace.create(name='KD_demo_workspace', subscription_id='da7f92e5-f637-4517-90fb-493d47170db2', resource_group='KD_resource_group'), name=diabetic_model_1, id=diabetic_model_1:2, version=2, tags={'data-format': 'CSV'}, properties={})
```

The screenshot shows the Azure AI | Machine Learning Studio interface. On the left, there's a sidebar with navigation links like Home, Model catalog (PREVIEW), and Models. The main area displays a model named 'diabetic_model_1:1'. The 'Details' tab is selected, showing the following attributes:

- Name:** diabetic_model_1
- Version:** 1
- Created on:** Jul 8, 2023 4:02 PM
- Created by:** Krishnendu Dey
- Type:** CUSTOM
- Created by job:** --
- Asset ID:** azureml://locations/centralindia/workspaces/8f813250-e78c-4309-9f63-32f488470208/models/diabetic_model_1/versions/1

On the right side, there are sections for Tags, Properties, and Description, each with a 'No [something]' message.

```
In [71]: ## Retrieving registered model
for model in Model.list(ws):
    print("Model Name : ", model.name, "|", "version :", model.version)
```

Model Name : diabetic_model_1 | version : 2
 Model Name : diabetic_model_1 | version : 1

```
In [ ]: ### An Alternate way of deploying model
# run.register_model( model_name='classification_model',
#                      model_path='outputs/diabeticmodel.pkl',           #
#                      description='A classification model',
#                      tags={'data-format': 'CSV'},
#                      model_framework=Model.Framework.SCIKITLEARN,
#                      model_framework_version='0.20.3')
```

STEP 20- Deploy Registered Model:

STEP 20.1: Getting the Registered ML Model

```
In [26]: ## Getting the Registered ML Model
model = Model(ws, model_name)           ## this is how a particular model is
model
```

```
Out[26]: Model(workspace=Workspace.create(name='KD_demo_workspace', subscription_id
='da7f92e5-f637-4517-90fb-493d47170db2', resource_group='KD_resource_grou
p'), name=diabetic_model_1, id=diabetic_model_1:1, version=1, tags={}, prop
erties={})
```

STEP 20.2 : Define an inference configuration

The model will be deployed as a service that consist of:

- A script to load the model and return predictions for submitted data.
- An environment in which the script will be run.

STEP 20.2.1 : Create an entry script

Create the entry script (sometimes referred to as the scoring script) for the service as a Python (.py) file. It must include two functions:

- init(): Called when the service is initialized.
- run(raw_data): Called when new data is submitted to the service. Typically, you use the init function to load the model from the model registry, and use the run function to generate predictions from the input data. The following example script of score.py file shows this pattern:

```
import json
import joblib
import numpy as np
import os
def init():
    global model
    # Get the path to the registered model file and
    # load it
    model_path = os.path.join(os.getenv('AZUREML_MODEL_DIR'), 'model.pkl')
    model = joblib.load(model_path)
```

STEP 20.2.2: Create an environment

Your service requires a Python environment in which to run the entry script, which you can define by creating an Environment that contains the required packages. This environment is required for inferencing the result of the ml model:

```
In [57]: ## Creating Inference Environment for the model prediction
from azureml.core.model import InferenceConfig
from azureml.core.conda_dependencies import CondaDependencies

myenv=Environment(name="demo-env")
conda_packages = ['numpy']
pip_packages = ['azureml-sdk','azureml-defaults','scikit-learn','pandas', 'n
mycondaenv = CondaDependencies.create(conda_packages=conda_packages,
                                         pip_packages=pip_packages,
                                         python_version='3.10.9')
myenv.python.conda_dependencies=mycondaenv
myenv.register(workspace=ws)
```



```
Out[57]: {  
    "assetId": "azureml://locations/centralindia/workspaces/8f813250-e78c-4  
309-9f63-32f488470208/environments/demo-env/versions/4",  
    "databricks": {  
        "eggLibraries": [],  
        "jarLibraries": [],  
        "mavenLibraries": [],  
        "pypiLibraries": [],  
        "rcranLibraries": []  
    },  
    "docker": {  
        "arguments": [],  
        "baseDockerfile": null,  
        "baseImage": "mcr.microsoft.com/azureml/openmpi4.1.0-ubuntu20.04:20  
230120.v1",  
        "baseImageRegistry": {  
            "address": null,  
            "password": null,  
            "registryIdentity": null,  
            "username": null  
        },  
        "buildContext": null,  
        "enabled": false,  
        "platform": {  
            "architecture": "amd64",  
            "os": "Linux"  
        },  
        "sharedVolumes": true,  
        "shmSize": null  
    },  
    "environmentVariables": {  
        "EXAMPLE_ENV_VAR": "EXAMPLE_VALUE"  
    },  
    "inferencingStackVersion": null,  
    "name": "demo-env",  
    "python": {  
        "baseCondaEnvironment": null,  
        "condaDependencies": {  
            "channels": [  
                "anaconda",  
                "conda-forge"  
            ],  
            "dependencies": [  
                "python=3.10.9",  
                {  
                    "pip": [  
                        "azureml-sdk~=1.49.0",  
                        "azureml-defaults~=1.49.0",  
                        "scikit-learn",  
                        "pandas",  
                        "numpy",  
                        "os"  
                    ]  
                },  
                "numpy"  
            ],  
            "name": "project_environment"  
        }  
    }  
}
```

```

        },
        "condaDependenciesFile": null,
        "interpreterPath": "python",
        "userManagedDependencies": false
    },
    "r": null,
    "spark": {
        "packages": [],
        "precachePackages": true,
        "repositories": []
    },
    "version": "4"
}

```

In [58]: # Combine the script and environment in an InferenceConfig

```

inference_config = InferenceConfig(entry_script='score.py',
                                   source_directory='.',
                                   environment=myenv)
inference_config

```

Out[58]: InferenceConfig(entry_script=score.py, runtime=None, conda_file=None, extra_docker_file_steps=None, source_directory=C:\Users\krish\OneDrive\Desktop\Study\Azure MLOPs\DP 100 Certifications, enable_gpu=None, base_image=None, base_image_registry=<azureml.core.container_registry.ContainerRegistry object at 0x0000027887695660>)

STEP 20.3 : Define a deployment configuration

Now that you have the entry script and environment, you need to configure the compute to which the service will be deployed. If you are deploying to an AKS cluster, you must create the cluster and a compute target for it before deploying:

In [59]: ### Specifying the machines of what configuration would be required to make ,

```

from azureml.core.webservice import AciWebservice
aciconfig = AciWebservice.deploy_configuration(cpu_cores=1, memory_gb=2)
aciconfig

```

Out[59]: <azureml.core.webservice.aci.AciServiceDeploymentConfiguration at 0x27885b58670>

STEP 20.4: Deployment

In [91]:

```
%time
## Deploying model as a service for real time predictions
service = Model.deploy(ws, "model-endpoint",
                       models=[model],
                       inference_config=inference_config,
                       deployment_config=aciconfig,
                       #deployment_target=predenv,
                       overwrite=True)
service.wait_for_deployment(show_output=True)
url = service.scoring_uri
print(url)
```

Tips: You can try `get_logs()`: <https://aka.ms/debugimage#dockerlog> (<https://aka.ms/debugimage#dockerlog>) or local deployment: <https://aka.ms/debugimage#debug-locally> (<https://aka.ms/debugimage#debug-locally>) to debug if deployment takes longer than 10 minutes.

Running

2023-07-08 19:18:25+05:30 Creating Container Registry if not exists.

2023-07-08 19:18:25+05:30 Registering the environment.

2023-07-08 19:18:25+05:30 Building image.

Failed

Service deployment polling reached non-successful terminal state, current service state: Unhealthy

Operation ID: 2c6d382b-6cf4-4289-9ed2-aeb40cb36fd6

More information can be found here:

Error:

```
{
  "code": "EnvironmentBuildFailed",
  "statusCode": 400,
  "message": "Failed building the Environment. You can try debugging locally first. Please refer to https://aka.ms/debugimage#debug-locally (https://aka.ms/debugimage#debug-locally) for more information."
}
```

```
-----  
WebserviceException                                     Traceback (most recent call last)  
File <timed exec>:8  
  
  File c:\Users\krish\anaconda3\envs\azure_ml_env\lib\site-packages\azureml\c  
ore\webservice\webservice.py:918, in Webservice.wait_for_deployment(self, s  
how_output, timeout_sec)  
    915         if not logs_response:  
    916             logs_response = 'Current sub-operation type not known,  
more logs unavailable.'  
--> 918         raise WebserviceException('Service deployment polling reach  
ed non-successful terminal state, current '  
    919                     'service state: {}\n'  
    920                     'Operation ID: {}\n'  
    921                     '{}\n'  
    922                     'Error:\n'  
    923                     '{}'.format(self.state, self._ope  
ration_endpoint.split('/')[-1],  
    924                                         logs_response, format  
_error_response), logger=module_logger)  
    925         print('{} service creation operation finished, operation "{}".  
format(self._webservice_type,  
    926  
operation_state))  
    927 except WebserviceException as e:  
  
WebserviceException: WebserviceException:  
    Message: Service deployment polling reached non-successful terminal  
state, current service state: Unhealthy  
Operation ID: 2c6d382b-6cf4-4289-9ed2-aeb40cb36fd6  
More information can be found here:  
Error:  
{  
    "code": "EnvironmentBuildFailed",  
    "statusCode": 400,  
    "message": "Failed building the Environment. You can try debugging locall  
y first. Please refer to https://aka.ms/debugimage#debug-locally (https://aka.ms/debugimage#debug-locally) for more information."  
}  
    InnerException None  
    ErrorResponse  
{  
    "error": {  
        "message": "Service deployment polling reached non-successful termi  
nal state, current service state: Unhealthy\nOperation ID: 2c6d382b-6cf4-42  
89-9ed2-aeb40cb36fd6\nMore information can be found here: \nError:\n{\n    \"code\": \"EnvironmentBuildFailed\",  
    \"statusCode\": 400,  
    \"message\": \"Failed building the Environment. You can try debugging locally first.  
Please refer to https://aka.ms/debugimage#debug-locally (https://aka.ms/debugimage#debug-locally) for more information.\n\""  
    }  
}
```

```
In [61]: print(service.get_logs())
```

```
Error in environment creation, more details may be found here: https://kdde  
mowostorage3c2969387.blob.core.windows.net/azureml/ImageLogs/e0e91736-ed72-  
4e35-b7a0-b3da81099b90/build.log (https://kddebowostorage3c2969387.blob.cor  
e.windows.net/azureml/ImageLogs/e0e91736-ed72-4e35-b7a0-b3da81099b90/build.  
log)
```

```
In [55]: ## Fetching the deployed service and examining it's health  
model_endpoint_name = "model-endpoint"  
service = AciWebservice(name = model_endpoint_name, workspace=ws)  
print(service.state)
```

Unhealthy

After Deployment is done, go to endpoint to see the serivce that has been deployed.

image.png

image.png

```
In [ ]: ## to get the logs if the service health is not "Healthy"  
# print(service.get_logs())
```

STEP 21: Testing the Deployed Model

```
In [63]: # Prepare your test data  
import pandas as pd  
import json  
data = pd.read_csv('test.csv')  
xt = data.values.tolist()  
test = json.dumps({"data":xt})
```

```
In [64]: test[:59]
```

```
Out[64]: '{"data": [[6.0, 148.0, 72.0, 35.0, 0.0, 33.6, 0.627, 50.0],'
```

```
In [ ]: # Fetching the deployed model
from azureml.core import Webservice
service = Webservice(workspace=ws, name="demo-model")
print(service.scoring_uri)
print(service.swagger_uri)
```

```
http://905ed832-9412-4764-886c-7b84317f6e2a.centralindia.azurecontainer.io/
score (http://905ed832-9412-4764-886c-7b84317f6e2a.centralindia.azurecontai
ner.io(score)
http://905ed832-9412-4764-886c-7b84317f6e2a.centralindia.azurecontainer.io/
swagger.json (http://905ed832-9412-4764-886c-7b84317f6e2a.centralindia.azur
econtainer.io/swagger.json)
```

```
In [ ]: # Running the predictions
prediction = service.run(input_data=test)
prediction[:10]
```

```
Out[56]: [1, 0, 0, 0, 1, 0, 1, 0, 1, 1]
```

```
In [ ]: ## Real-Time predictions using REST call
import requests
headers = {'content-Type': 'application/json'}

response = requests.post(service.scoring_uri, test, headers=headers)
response.text
```

NOTE: Local Deployment in Docker

```
In [68]: ## It will work when Docker is installed in the PC
from azureml.core.webservice import LocalWebservice
deployment_config = LocalWebservice.deploy_configuration(port=5002)
local_service = Model.deploy(ws,
                             'test-deployment-1',
                             [model],
                             inference_config=inference_config,
                             deployment_config=deployment_config)
```

Failed to create Docker client. Is Docker running/installed?
When you deploy locally, we download a dockerfile
execute docker build on it, and docker run the built container for you
Error: Error while fetching server API version: (2, 'CreateFile', 'The
system cannot find the file specified.')

```
-----
error                                         Traceback (most recent call
last)
File c:\Users\krish\anaconda3\envs\azure_ml_env\lib\site-packages\dock
er\api\client.py:214, in APIClient._retrieve_server_version(self)
    213     try:
--> 214         return self.version(api_version=False)["ApiVersion"]
    215     except KeyError:

File c:\Users\krish\anaconda3\envs\azure_ml_env\lib\site-packages\dock
er\api\daemon.py:181, in DaemonApiMixin.version(self, api_version)
    180     url = self.url + "/version"  versioned api=api version)
```

In []:

In []:

In []:

In []:

STEP 22:Authentication

In production, you will likely want to restrict access to your services by applying authentication. There are two kinds of authentication you can use:

- **Key:** Requests are authenticated by specifying the key associated with the service.
- **Token:** Requests are authenticated by providing a JSON Web Token (JWT).

By default, authentication is disabled for ACI services, and set to key-based authentication for AKS services (for which primary and secondary keys are automatically generated). You can optionally configure an AKS service to use token-based authentication (which is not supported for ACI services).

Assuming you have an authenticated session established with the workspace, you can retrieve the keys for a service by using the get_keys method of the WebService object associated with the service:

```
In [ ]: # primary_key, secondary_key = service.get_keys()  
# print(primary_key)
```

```
In [ ]: # Set additional parameter key_or_token to make predictions from the deployed  
import requests  
headers = { "Content-Type": "application/json",  
            "Authorization": "Bearer " + key_or_token }  
  
response = requests.post(service.scoring_uri, test, headers=headers)  
response.text
```